B.Maheshwar(192411195)

Day4-Data Mining Experiments

1. Customer Segmentation using Clustering

Aim:

To segment supermarket customers using clustering based on income and spending score.

Algorithm:

1. Collect customer data

2. Select income and spending score

3. Apply K-Means

4. Choose k clusters

5. Analyze clusters

Output:

Clusterer output

```
Initial starting points (random):

Cluster 0: 4,23,Female,16,77
Cluster 1: 2,21,Male,15,81
Cluster 2: 1,19,Male,15,39
Cluster 3: 3,20,Female,16,6
Cluster 4: 5,31,Female,17,40

Missing values globally replaced with mean/mode

Final cluster centroids:
                                Cluster#
Attribute          Full Data          0          1          2          3          4
                     (5.0)       (1.0)      (1.0)      (1.0)      (1.0)      (1.0)
================================================================================
CustomerID                3          4          2          1          3          5
Age                    22.8         23         21         19         20         31
Gender               Female     Female       Male       Male     Female     Female
AnnualIncome           15.8         16         15         15         16         17
SpendingScore          48.6         77         81         39          6         40



Time taken to build model (full training data) : 0 seconds

=== Model and evaluation on training set ===

Clustered Instances

0       1 ( 20%)
1       1 ( 20%)
2       1 ( 20%)
3       1 ( 20%)
4       1 ( 20%)
```

2. Employee Data Clustering using K-Means

Aim:

To perform clustering on employee data using K-Means.

Algorithm:

1. Load CSV file

2. Convert to ARFF

3. Apply K-Means

4. Change cluster size

5. Visualize output

Output:

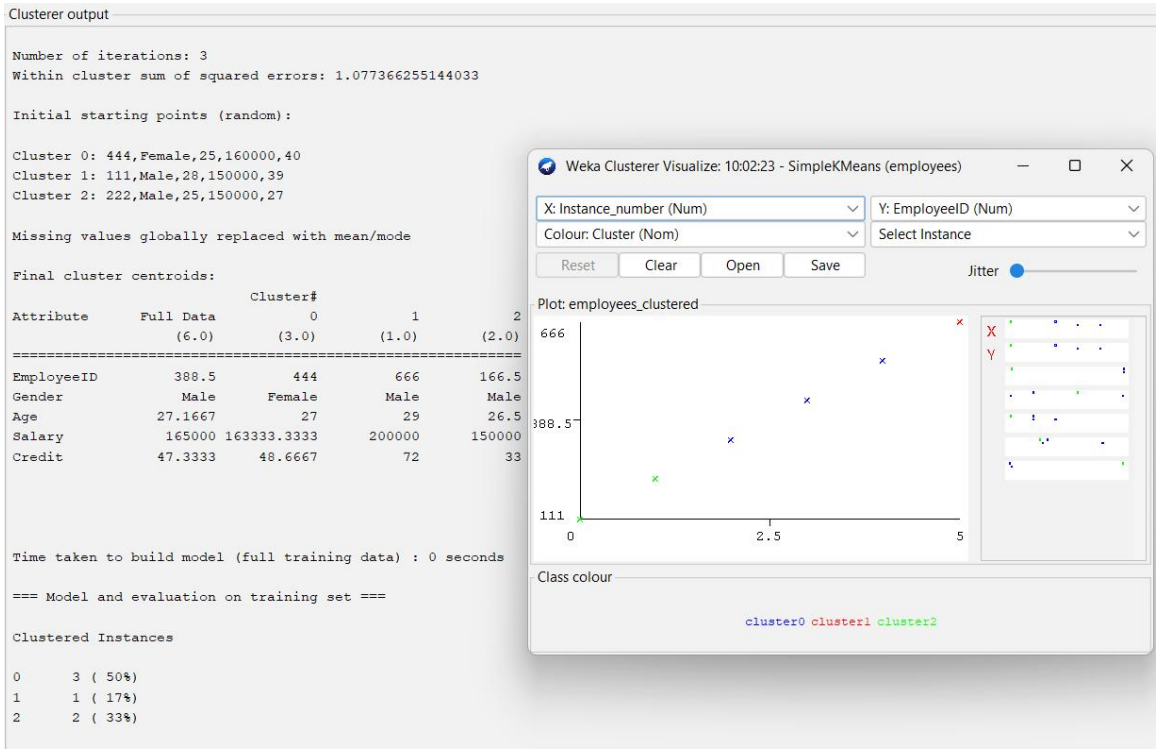

3. Naive Bayes Classification

Aim:

To classify categorical data using Naive Bayes and compare with SVM.

Algorithm:

1. Load dataset

2. Apply Naive Bayes

3. Apply SVM

4. Compare accuracy

Output:

Classifier output

```
Instances:    6
Attributes:   5
              outlook
              temperature
              humidity
              windy
              play
Test mode:    10-fold cross-validation

=== Classifier model (full training set) ===

SMO

Kernel used:
  Linear Kernel: K(x,y) = <x,y>

Classifier for classes: yes, no

BinarySMO

Machine linear: showing attribute weights, not support vectors.

          1       * (normalized) outlook=sunny
 +       -0.7     * (normalized) outlook=overcast
 +       -0.3     * (normalized) outlook=rainy
 +        0.145   * (normalized) temperature
 +       -0.3231  * (normalized) humidity
 +       -1       * (normalized) windy=false
 +        0.9435

Number of kernel evaluations: 21 (79% cached)



Time taken to build model: 0.01 seconds
```

4. Vegetarian Count

Aim:

To find number of vegetarian and non-vegetarian persons.

Algorithm:

1. Load data

2. Count yes and no

3. Compare totals

Output:

Selected attribute

Name: Vegetarian Type: Nominal
Missing: 0 (0%) Distinct: 2 Unique: 0 (0%)

No.	Label	Count	Weight
1	yes	7	7
2	no	3	3

Class: Vegetarian (Nom) Visualize All

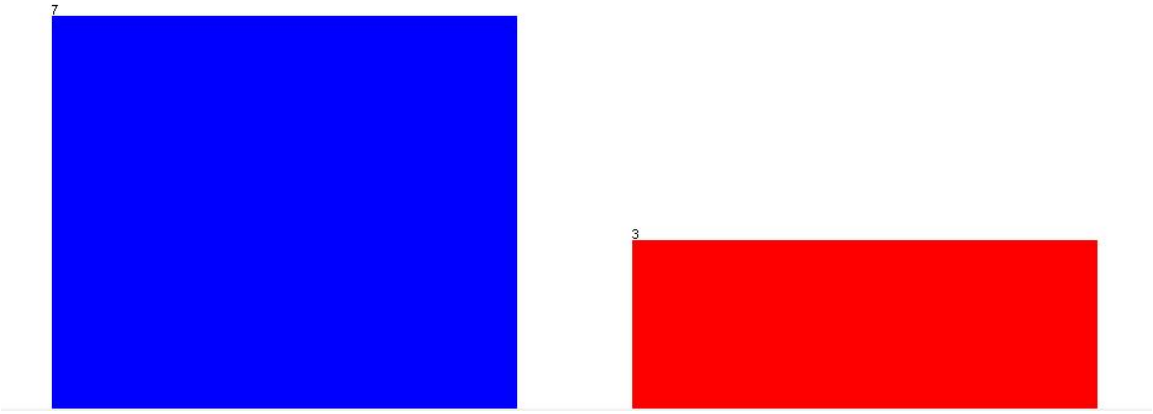

5. Scatter Plot

Aim:

To plot scatter graph for mobile sales data.

Algorithm:

1. Load x,y data

2. Choose scatter plot

3. Plot graph

Output:

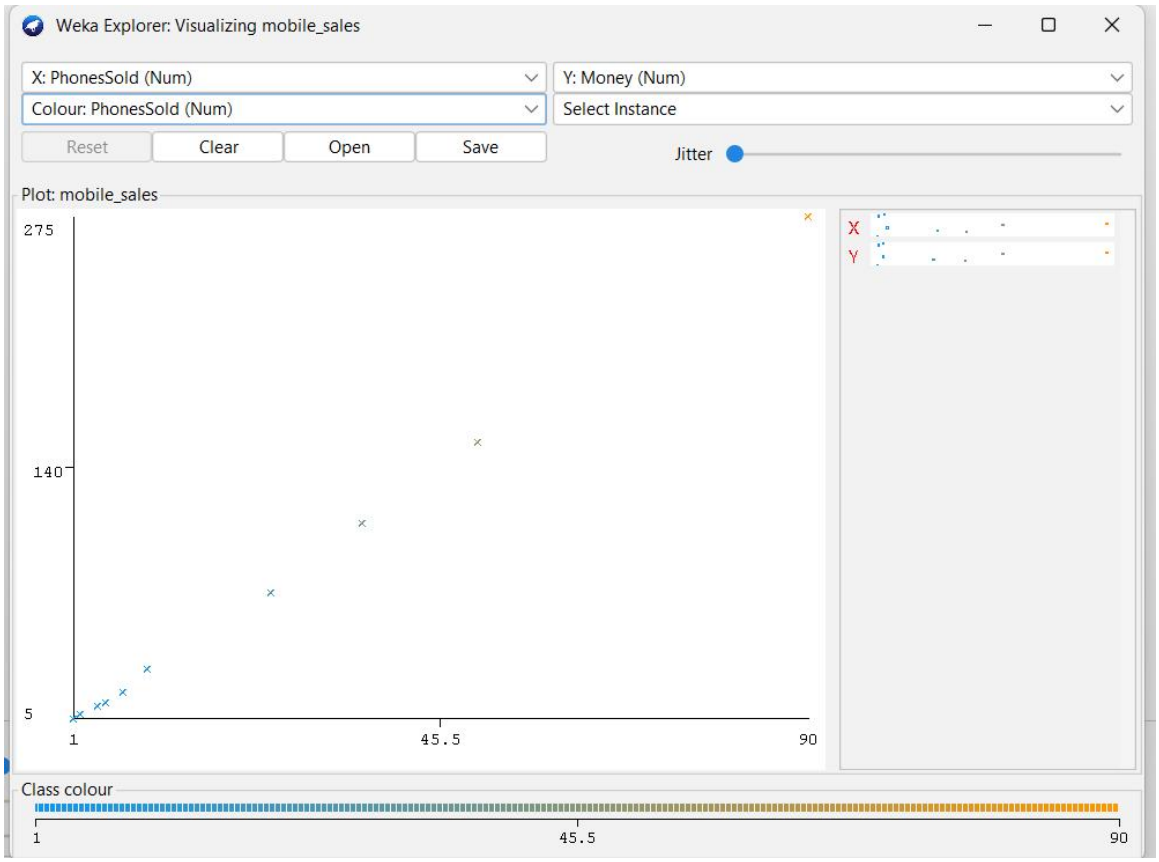

6. FP-Growth Algorithm

Aim:

To generate association rules using FP-Growth.

Algorithm:

1. Load transactions

2. Set min support

3. Build FP-tree

4. Generate rules

Output:

7. Diabetes Prediction

Aim:

To predict diabetes using Decision Tree and compare with SVM.

Algorithm:

1. Load dataset

2. Apply Decision Tree

3. Apply SVM

4. Compare accuracy and F1

Output:

Classifier output

```
Scheme:       weka.classifiers.functions.SMO -C 1.0 -L 0.001 -P 1.0E-12 -N 0 -V -1 -W 1 -K "w
Relation:     diabetes
Instances:    5
Attributes:   5
              Glucose
              BloodPressure
              BMI
              Age
              Class
Test mode:    10-fold cross-validation

=== Classifier model (full training set) ===

SMO

Kernel used:
  Linear Kernel: K(x,y) = <x,y>

Classifier for classes: Positive, Negative

BinarySMO

Machine linear: showing attribute weights, not support vectors.

        -1.5338 * (normalized) Glucose
 +       0.4149 * (normalized) BloodPressure
 +      -0.3828 * (normalized) BMI
 +      -0.7273 * (normalized) Age
 +       0.4981

Number of kernel evaluations: 12 (74.468% cached)



Time taken to build model: 0.01 seconds
```

8. Data Binning

Aim:

To partition data using equal-width, equal-frequency and clustering.

Algorithm:

1. Load marks

2. Apply binning methods

3. Plot histogram

Output:

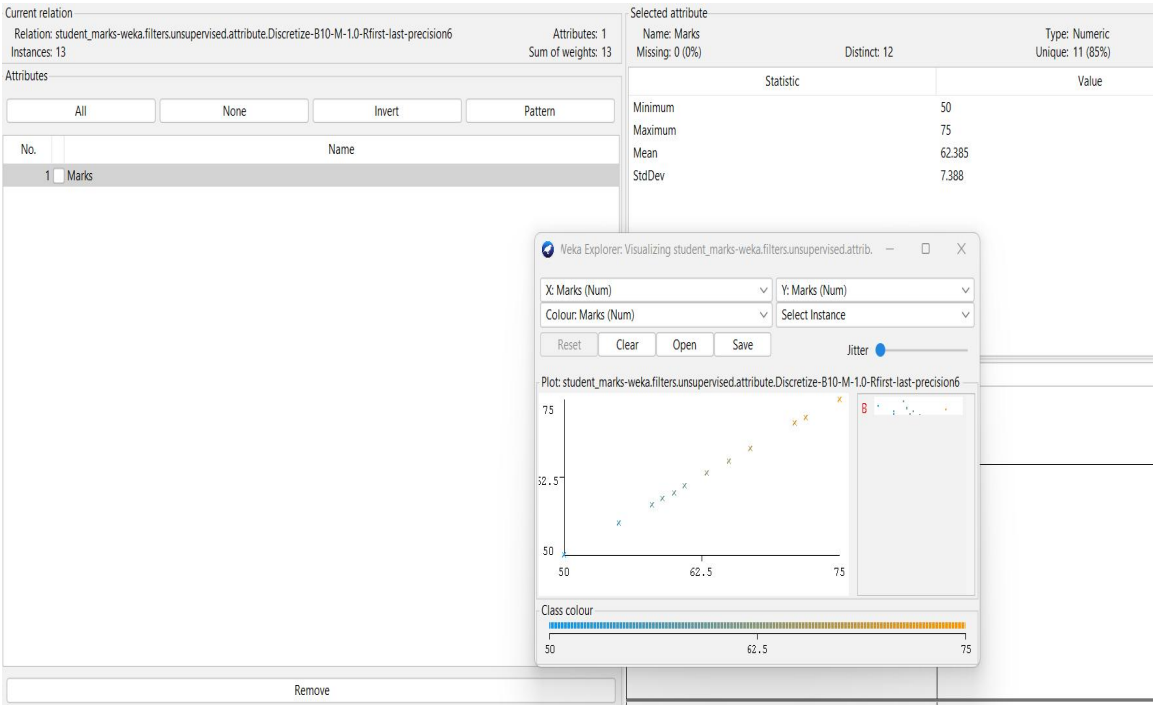

9. Decision Tree Dataset

Aim:

To create dataset and generate rules from decision tree.

Algorithm:

1. Create ARFF

2. Build tree

3. Generate rules

4. Plot confusion matrix

Output:

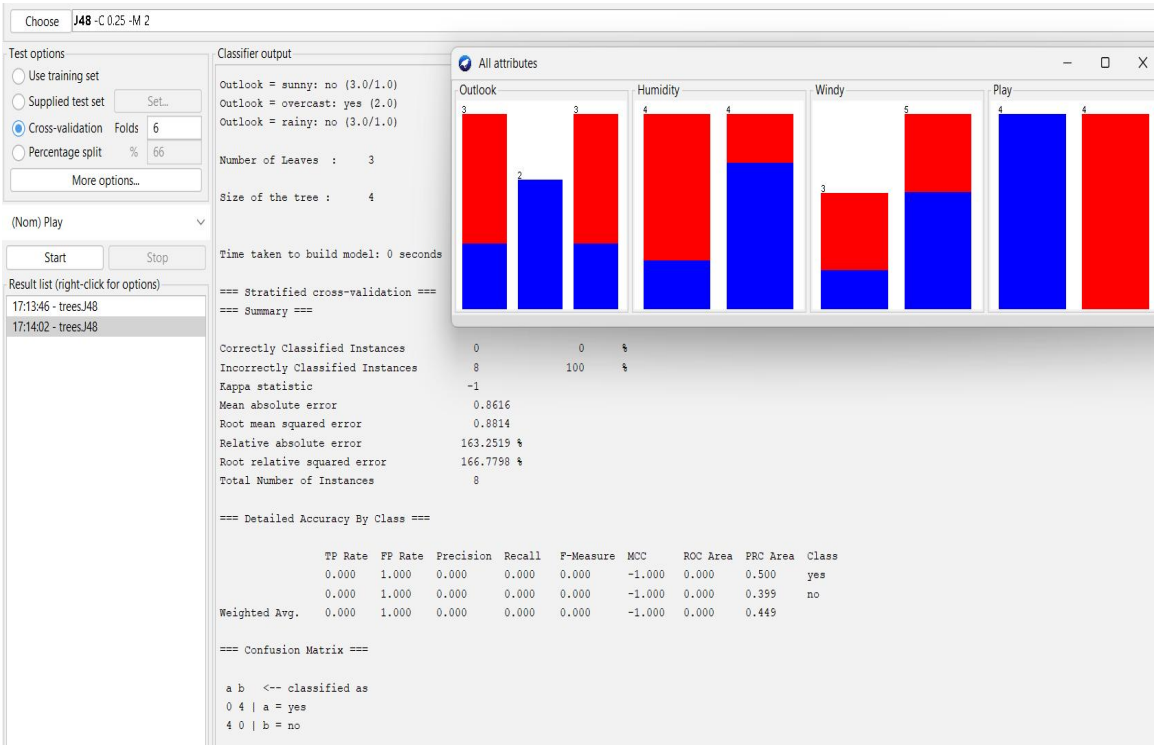

10. Apriori vs FP-Growth

Aim:

To compare Apriori and FP-Growth algorithms.

Algorithm:

1. Load transaction data

2. Apply Apriori

3. Apply FP-Growth

4. Compare rules

Output:

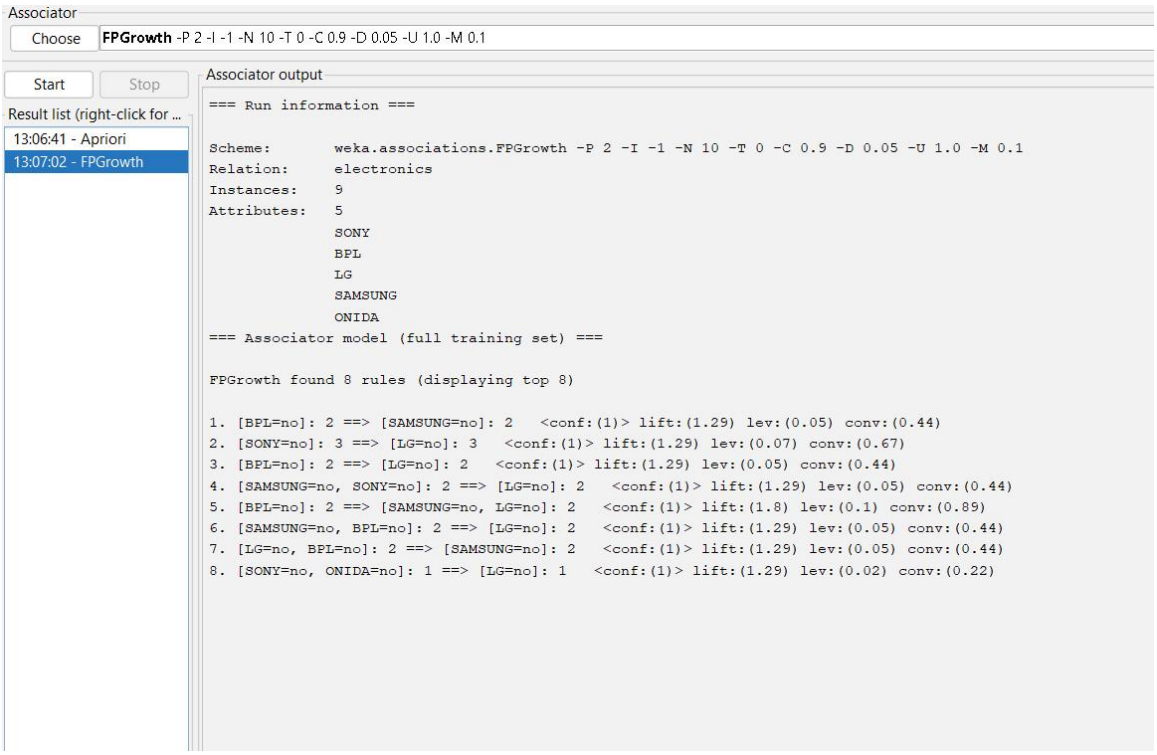

11. Normalization Techniques

Aim:

To normalize data using different normalization methods.

Algorithm:

1. Load data

2. Apply min-max

3. Apply z-score

4. Apply decimal scaling

Output:

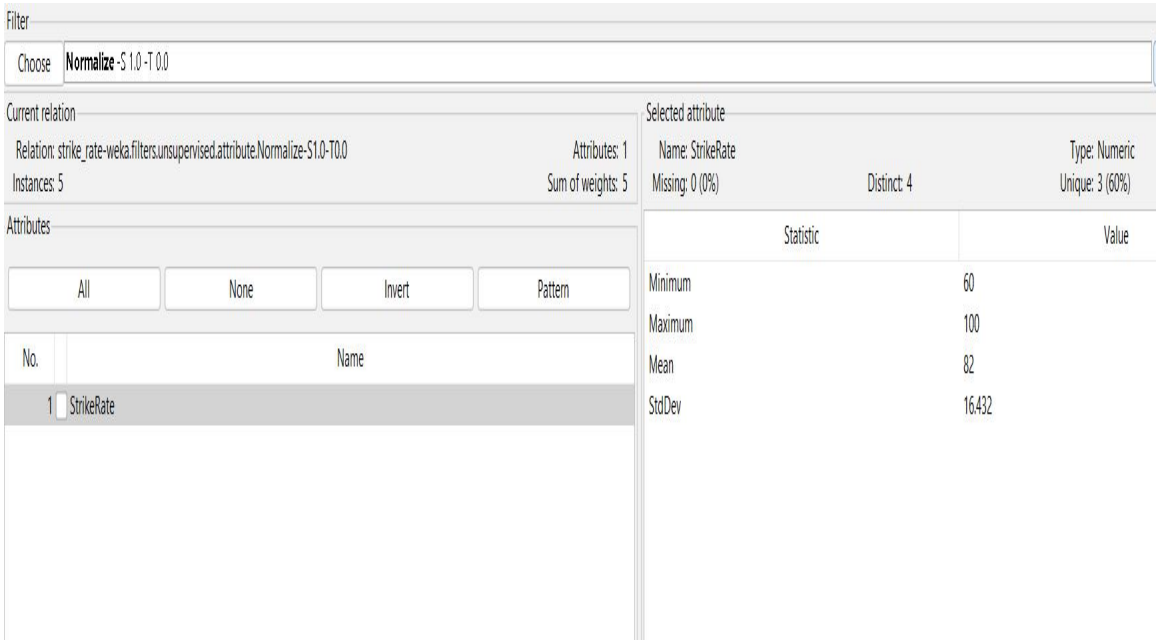

12. Mean and Variance Calculation

Aim:

To calculate variance and standard deviation for given data.

Algorithm:

1. Load data

2. Compute mean

3. Compute variance

4. Compute standard deviation

Output:

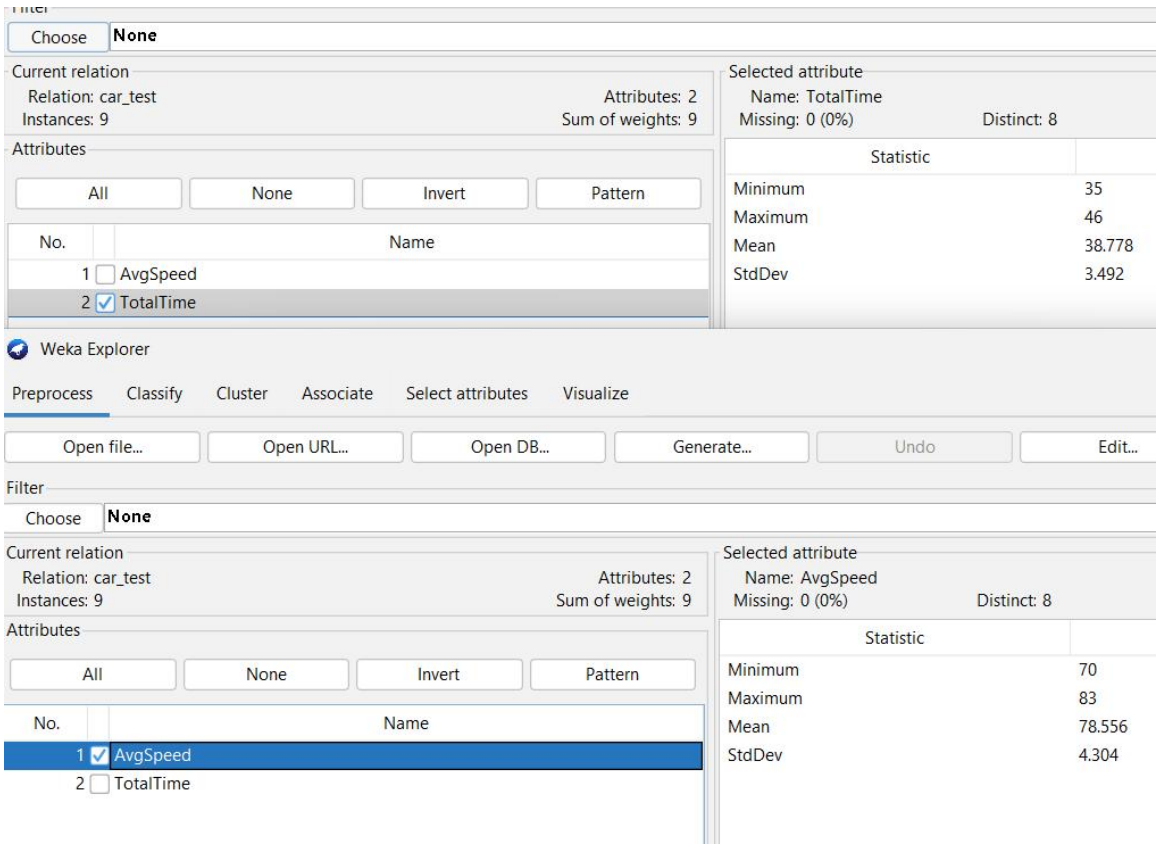

13. Frequent Itemset Mining (Apriori & FP-Growth)

Aim:

To find frequent itemsets and generate association rules using Apriori and FP-Growth.

Algorithm:

1. Load transaction dataset

2. Set minimum support and confidence

3. Apply Apriori

4. Apply FP-Growth

5. Compare frequent itemsets and rules

Output:

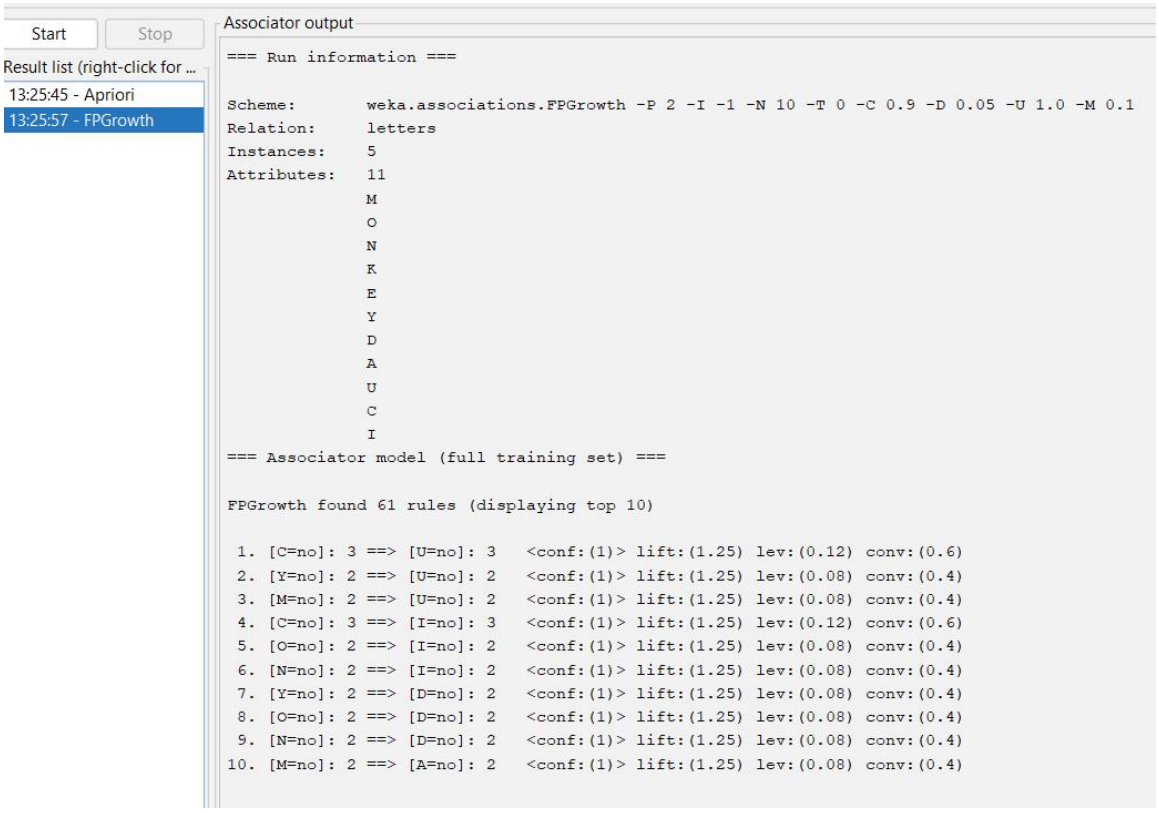

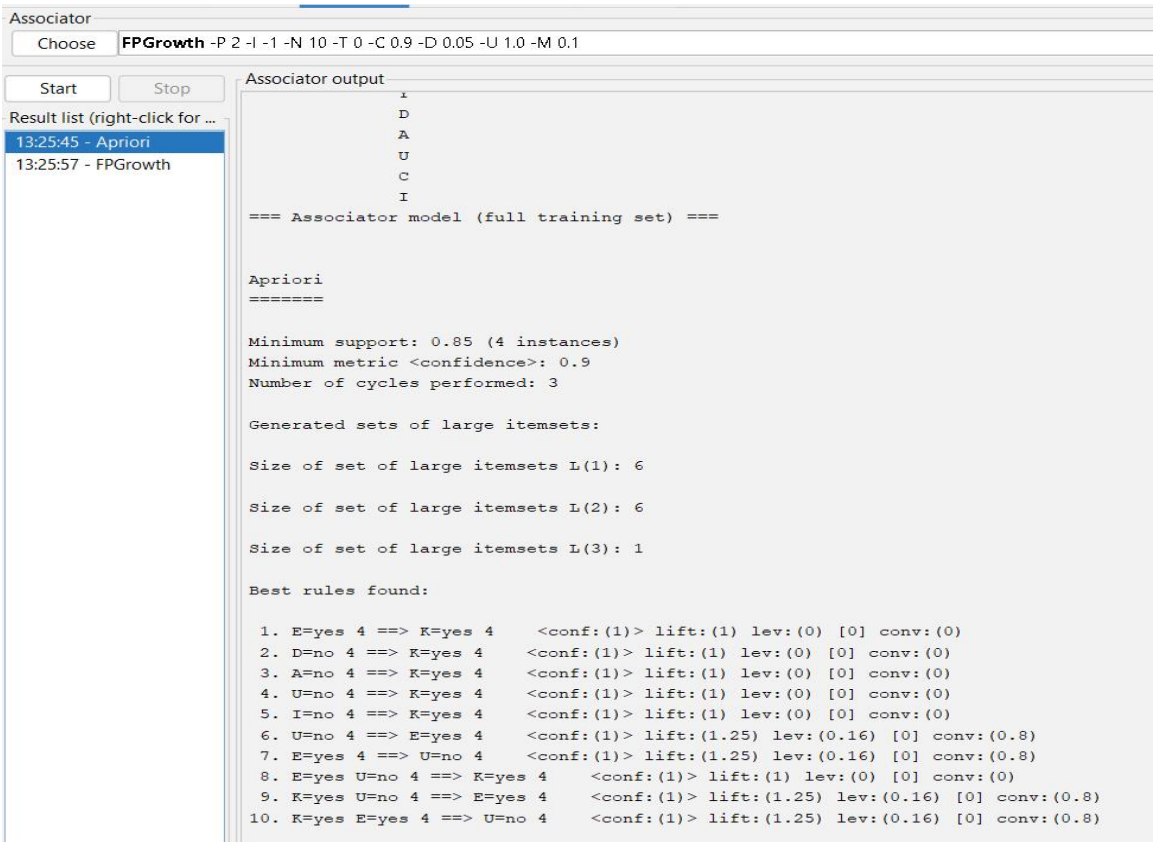
Associator
Choose
FPGrowth -P 2 -I -1 -N 10 -T 0 -C 0.9 -D 0.05 -U 1.0 -M 0.1
Start
Stop
Result list (right-click for ...
13:25:45 - Apriori
13:25:57 - FPGrowth
Associator output
I
D
A
U
C
I
=== Associator model (full training set) ===
Apriori
=======
Minimum support: 0.85 (4 instances)
Minimum metric <confidence>: 0.9
Number of cycles performed: 3
Generated sets of large itemsets:
Size of set of large itemsets L(1): 6
Size of set of large itemsets L(2): 6
Size of set of large itemsets L(3): 1
Best rules found:
1. E=yes 4 ==> K=yes 4 <conf:(1)> lift:(1) lev:(0) [0] conv:(0)
2. D=no 4 ==> K=yes 4 <conf:(1)> lift:(1) lev:(0) [0] conv:(0)
3. A=no 4 ==> K=yes 4 <conf:(1)> lift:(1) lev:(0) [0] conv:(0)
4. U=no 4 ==> K=yes 4 <conf:(1)> lift:(1) lev:(0) [0] conv:(0)
5. I=no 4 ==> K=yes 4 <conf:(1)> lift:(1) lev:(0) [0] conv:(0)
6. U=no 4 ==> E=yes 4 <conf:(1)> lift:(1.25) lev:(0.16) [0] conv:(0.8)
7. E=yes 4 ==> U=no 4 <conf:(1)> lift:(1.25) lev:(0.16) [0] conv:(0.8)
8. E=yes U=no 4 ==> K=yes 4 <conf:(1)> lift:(1) lev:(0) [0] conv:(0)
9. K=yes U=no 4 ==> E=yes 4 <conf:(1)> lift:(1.25) lev:(0.16) [0] conv:(0.8)
10. K=yes E=yes 4 ==> U=no 4 <conf:(1)> lift:(1.25) lev:(0.16) [0] conv:(0.8)